

universidade de aveiro



deti

departamento de eletrónica,  
telecomunicações e informática

# Software Engineering

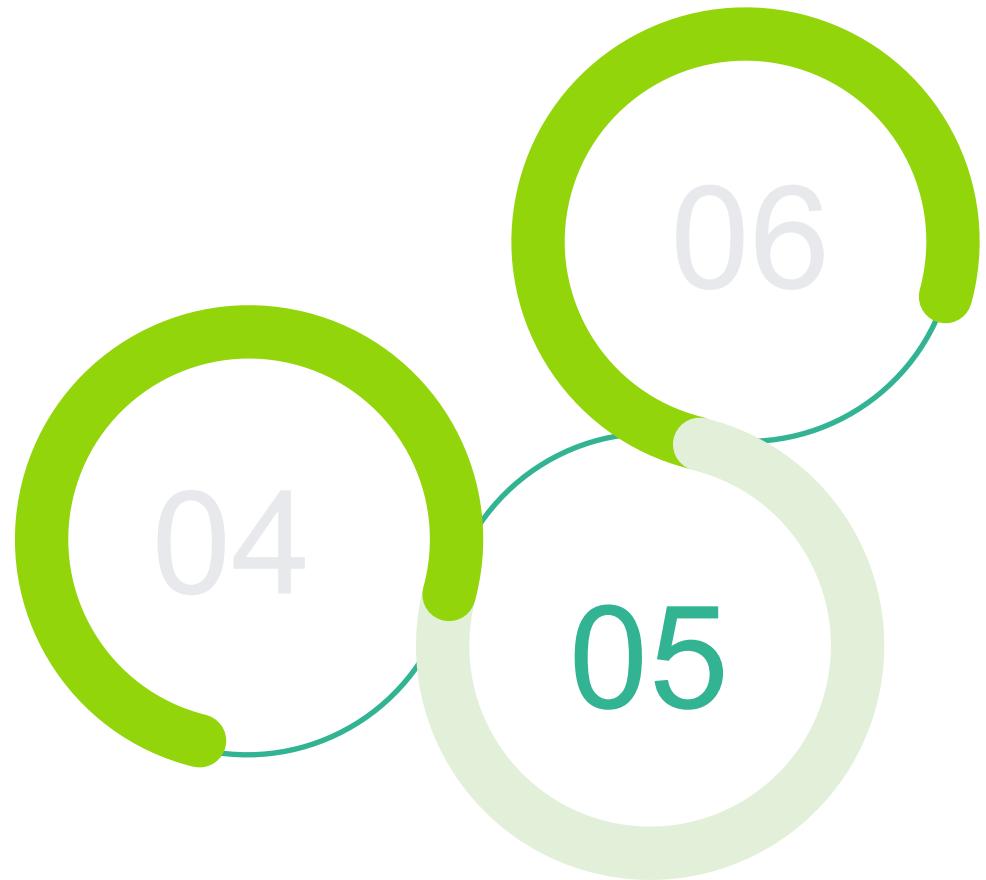
Strategies for continuous delivery

2025/2026



# Outline

- Core principles of software delivery
- Feature toggling
- Progressive delivery strategies
- Continuously releasing user-installed software



# Minimal Viable Product

- Scenario: The client provided ideas, and the team developed the core features.
- Current status:
  - Core functionality is **implemented and working**.
  - Product is **usable by early adopters** for testing and feedback.
  - Scope is **minimal**, focusing only on **essential features**.



Haven't we done  
this before?

# Release candidate

- A change may or may not be **releasable**
  - The code alone doesn't tell us.
- The build, deployment, and test process **validates whether a change can be safely released**.
  - This process increases confidence that the change is safe for production.
- Applies to any change: new functionality, bug fixes, or performance tuning.



# Core principles of software delivery

**Create a repeatable, reliable process for releasing software**

- The repeatability and reliability derive from two principles:
  - **Automate** almost everything
  - Keep everything you need to build, deploy, test, and release your application in **version control**
- Deploying software ultimately involves three things:
  - Provisioning and managing the environment in which your application will run
    - Hardware configuration, software, infrastructure, and external services
  - Installing the correct version of your application into it
  - Configuring the application, including any data or state it requires
- Application deployment can be implemented using a fully automated process from version control
  - Application configuration can also be a fully automated process, with the necessary scripts and state kept in version control or databases

# Core principles of software delivery

## Automate almost everything

- Some tasks **cannot** be automated
  - Exploratory testing, user demos, compliance approvals
- Most of the release process **can** be automated
  - Builds, deployments, acceptance tests, database migrations, network/firewall configs
- Automation reduces repeated manual effort and ensures consistency
- Start gradually
  - Automate bottlenecks first, expand over time
- Automation is essential for an efficient deployment pipeline
  - People get what they need at the push of a button

# Core principles of software delivery

## Keep everything in version control

- Keep everything needed to build, deploy, test, and release under version control:
  - Requirement documents, test scripts, automated tests
  - Deployment, database, network, and stack configuration scripts
  - Libraries, toolchains, and technical documentation
- Each build should have a single identifier linking all relevant files
  - e.g., build number or change set
- A new team member should be able to:
  - Check out the repository
  - Run a single command to build and deploy the application

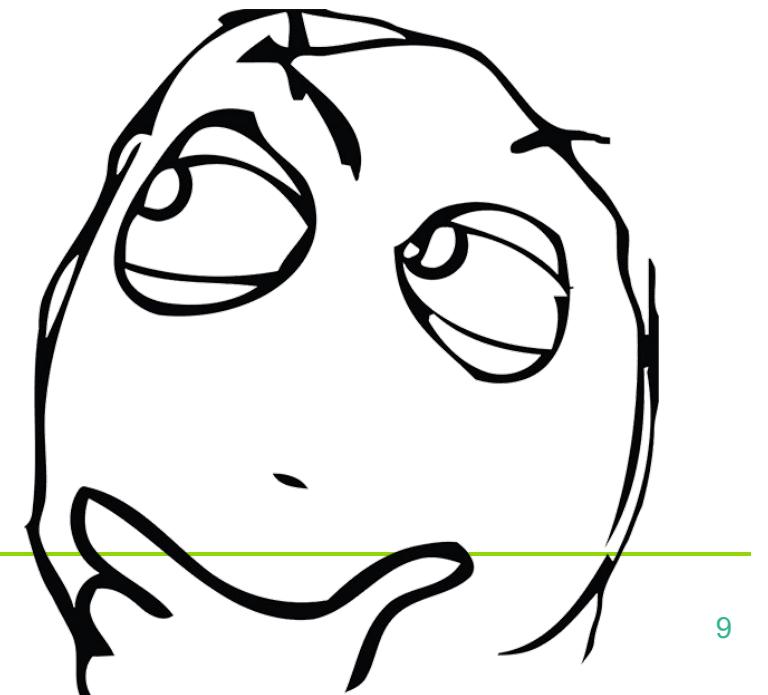
# Core principles of software delivery

## Done means released

- There is no “80% done.”
  - Things are either done, or they are not.
- It is possible to estimate the work remaining before something is done
  - But those will only ever be estimates.
- This principle has an interesting corollary
  - It is not in the power of one person to get something done.
  - It requires a number of people on a delivery team to work together to get anything done.
  - That's why it's so important for everybody to work together from the beginning.
- It's also why the whole delivery team is responsible for delivering

# Software delivery

- Same scenario: Client liked the product, now he wants more features.
  - How should we deliver new features to him?
  - Should we release everything at once or gradually?
  - Do we force updates, or give users a choice?
  - How do we minimize risk while keeping users happy?
  - What metrics or feedback should guide our release decisions?
  
- **Would you rather risk upsetting a few users with automatic updates or risk losing everyone by delaying new features?**



# Feature toggling

## Overview

- Feature toggling (also called feature flags) is a software development technique that allows features to be **turned on or off** without deploying new code.
  - Instead of releasing a new feature to all users immediately, it enables control over who can access it and when.



```
if feature_flag_enabled("new_dashboard"):  
    show_new_dashboard()  
else:  
    show_old_dashboard()
```

# Feature toggling

## Motivation

- Continuous delivery
  - Ship code to production even if features aren't ready.
- A/B testing & experimentation
  - Test features on a subset of users.
- Safe rollbacks
  - Disable a feature instantly if issues appear.
- Targeted releases
  - Roll out features to specific users, regions, or accounts.
- Decouple deployment from release
  - Code can be deployed independently of feature activation.



# Feature toggling

## Four main types

- **Release toggles**
  - Control **when a feature is visible**.
  - Temporary; removed after full release.
  - Example: Releasing a new payment system to 10% of users.
- **Experiment toggles**
  - For **A/B testing**.
  - Determine which variant of a feature performs better.
  - Example: Two versions of a signup form.
- **Ops toggles (Operational toggles)**
  - Control **operational aspects** rather than functionality.
  - Useful for **performance tuning, enabling/disabling features under load**.
  - Example: Toggle caching or logging level.
- **Permission toggles (or permanent toggles)**
  - Enable features for **specific users or groups**.
  - Example: Premium features enabled for paying customers.
  - Can be **long-lived**.

# Feature toggling

## Four main types

- **Release toggles**
  - Control **when a feature is visible**.
  - Temporary; removed after full release.
  - Example: Releasing a new payment system to 10% of users.
- **Experiment toggles**
  - For **A/B testing**.
  - Determine which variant of a feature performs better.
  - Example: Two versions of a signup form.
- **Ops toggles (Operational toggles)**
  - Control **operational aspects** rather than functionality.
  - Useful for **performance tuning, enabling/disabling features under load**.
  - Example: Toggle caching or logging level.
- **Permission toggles (or permanent toggles)**
  - Enable features for **specific users or groups**.
  - Example: Premium features enabled for paying customers.
  - Can be **long-lived**.

# Continuously delivery



# Progressive delivery strategies

## Eat you own dog food

- Idea: Dev team (and internal staff) use the newest version internally
- If everything is as expected it is roll it out to further users



# Progressive delivery strategies

## Canary releases

- Idea: release a new version to a subset of users first, while all other users interact with the old, stable version
  - In case of issues with new version, only a small number of users is affected
- The name comes from the “canary in a coal mine”
  - Miners used canaries to detect toxic gases before it was too late. If the canary showed signs of distress, the miners knew there was danger — the same principle applies here.
- User selection based on:
  - (geographic) location
  - role (admin, early, stable, etc.)
  - random basis (e.g., 1% of traffic)



# Progressive delivery strategies

## Canary releases

- Canary releases can be implemented at **different levels**:

- a) Infrastructure / Load Balancer Level

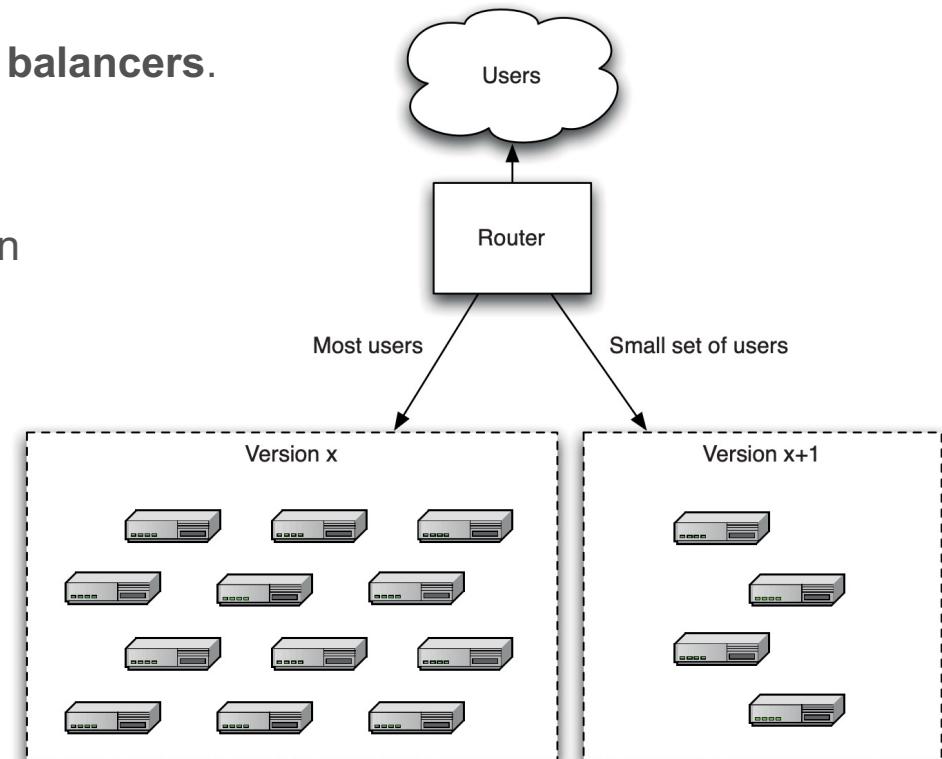
- Traffic is routed selectively using **reverse proxies or load balancers**.

- b) Application level

- The application decides which users get the canary version
      - (based on user ID, region, feature flag, etc.).
    - Often combined with **feature toggles**.

- c) CI/CD Pipeline Level

- Orchestrated directly in the pipeline.



# Progressive delivery strategies

## Canary releases

- Monitoring is *crucial* for canary releases. Key indicators include:
  - Error rates
  - Response times
  - CPU/memory usage
  - User satisfaction metrics
  - Business KPIs (e.g., conversion rate, retention)
- Automation can compare canary metrics to the baseline and **automatically decide whether to continue or rollback**
  - Known as Automated Canary Analysis (ACA)

# Progressive delivery strategies

## Canary releases

➤ Let's say version 2.0 of an API is ready.

- Step 1: Deploy it to **5% of users**.
- Step 2: Monitor **latency, error rate, CPU load**, etc.
- Step 3: If metrics are healthy → increase rollout to **20%, 50%, 100%**.
- Step 4: If metrics degrade → **rollback to version 1.9** for everyone.

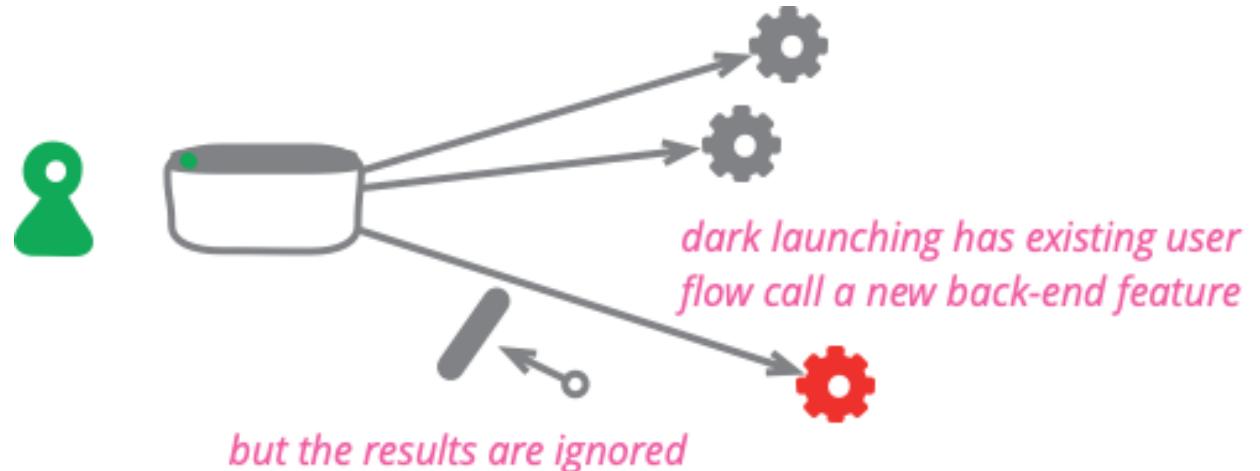
➤ Key Objectives

- Reduce risk of production failures.
- Validate functionality under real-world conditions.
- Gather early feedback from a small audience.
- Enable fast rollback if the release introduces issues.

# Progressive delivery strategies

## Dark launches

- A deployment strategy where a feature is released to production
- Feature is **not yet exposed** to end-users
- Used to test scalability, performance, and stability
- Often combined **with feature toggles** or **hidden endpoints**



*this allows us to monitor the load on existing systems before the feature is officially released*

# Progressive delivery strategies

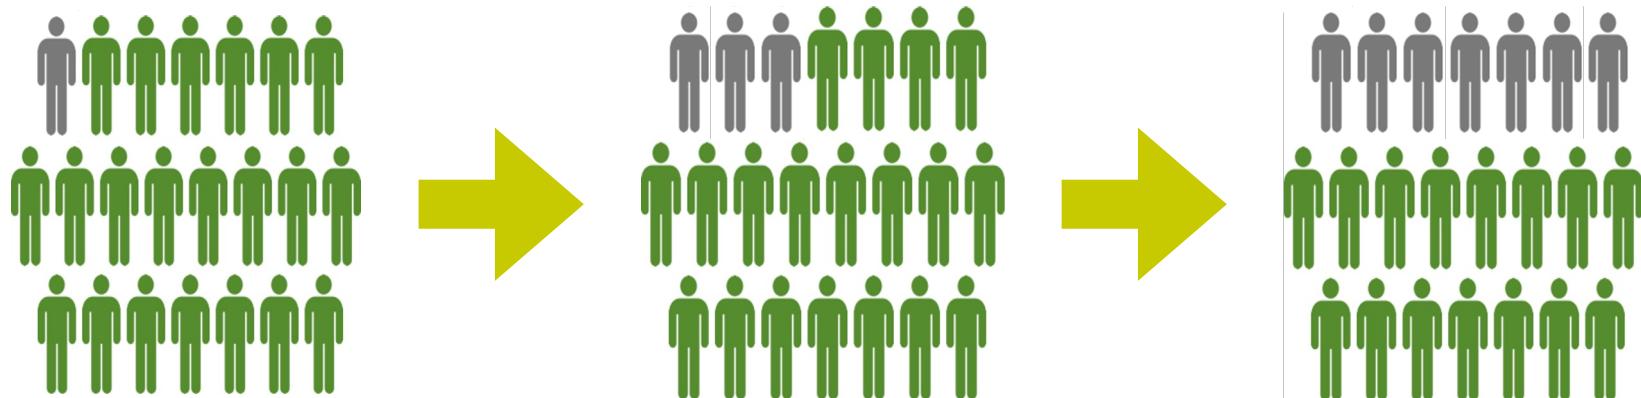
## Dark launches

- Identify **performance bottlenecks** before release
- Detect infrastructure or **integration issues**
- Validate **resource usage under real conditions**
- Ensure smooth user experience once the feature is public
- Core idea
  - Deploy new features **into production**
  - Keep them **invisible to users**
  - Observe how they behave under real production load

# Progressive delivery strategies

## Gradual rollouts

- Gradually increase the number of users assigned to the newest version
- Continue the rollout step by step until it fully replaces the previous version
- Allows early detection of issues before full release



# Progressive delivery strategies

## Gradual rollouts

### ➤ Workflow

- Start with a small subset of users
- Monitor performance and errors
- Gradually expand user percentage
- Once stable, complete the rollout

### ➤ Purpose

- To test scalability and stability under real-world conditions
- To observe system behaviour with increasing load
- To minimise risk of widespread failure

# Progressive delivery strategies

## A/B Testing

- A/B Testing compares two or more versions of a feature or interface
- Users are randomly divided into groups (A, B, etc.)
- Each group experiences a different version
- Results determine which version performs better

# Progressive delivery strategies

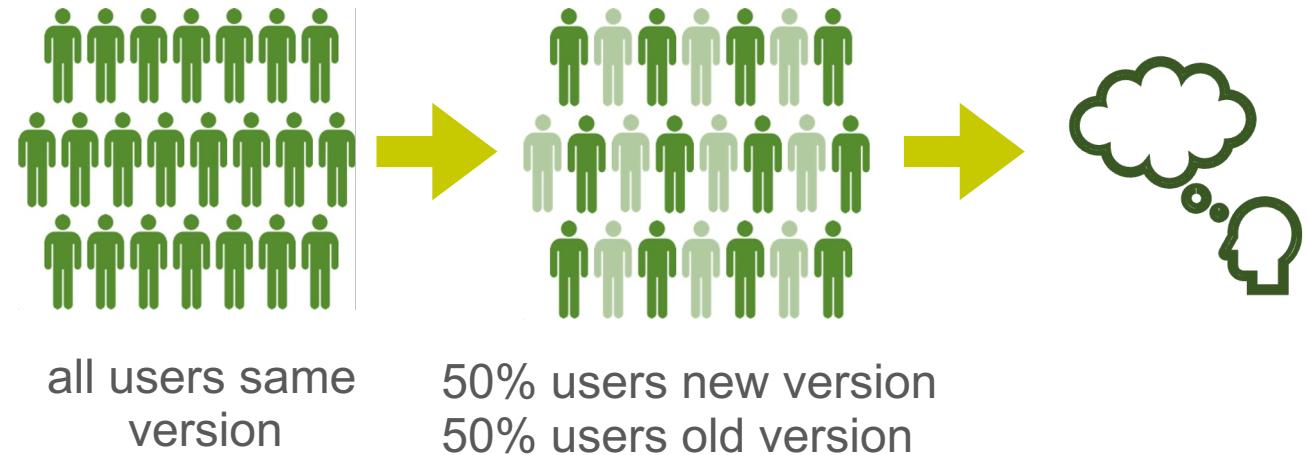
## A/B Testing

### ➤ Purpose

- To evaluate user behaviour and measure impact of design or functional changes
- To validate hypotheses before full rollout
- To support evidence-based decision-making

### ➤ Process Overview

- Define a goal or success metric
  - (e.g., click rate, engagement, conversion)
- Split users into test groups
- Deploy each variant
- Collect and analyse data statistically
- Decide which version to adopt or discard



# Progressive delivery strategies

## A/B Testing

- Can be used within a progressive rollout to **test user reactions**
- A/B Testing focuses **on comparison**, while progressive rollout focuses on scale and stability
- Combining both ensures functional and performance validation
- A/B Testing helps organisations:
  - Make data-backed decisions
  - Understand user preferences
  - Deliver better-performing features

# Progressive delivery strategies

## Blue/Green Deployments

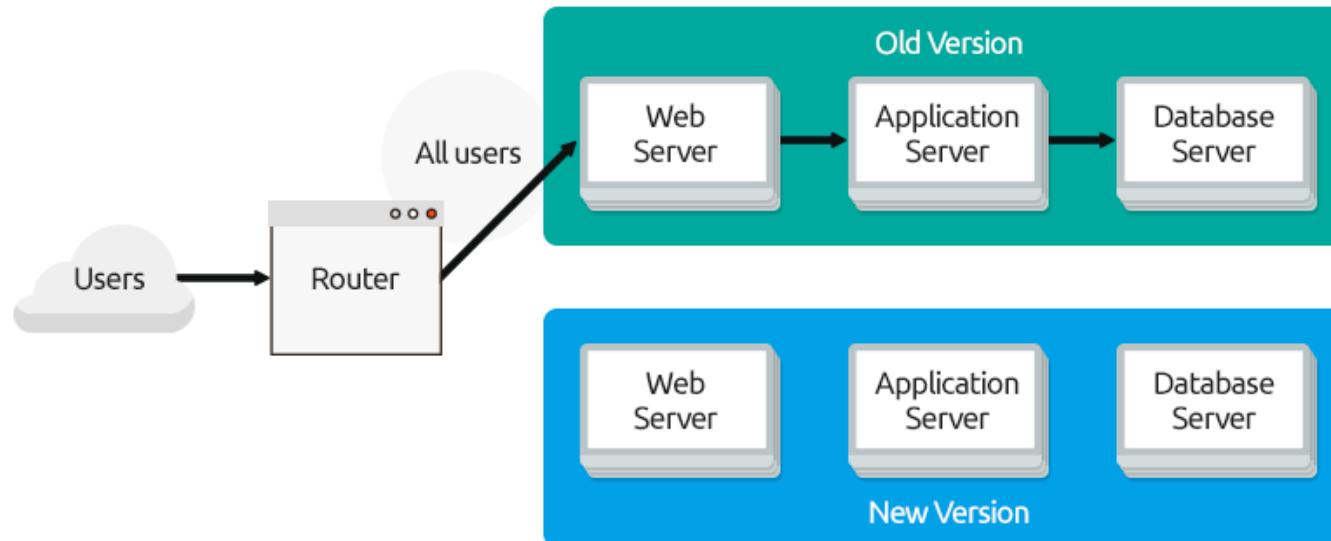
- Maintain two identical production environments:
  - Green: current, stable version
  - Blue: new version under preparation
- The router directs traffic to either environment
- Switching traffic effectively releases the new version

# Progressive delivery strategies

## Blue/Green Deployments

### ➤ Release process

1. Deploy the new version to the inactive environment (e.g., Blue)
2. Test the new version in isolation
3. Once validated, switch traffic from Green → Blue
4. Keep the previous environment ready for instant rollback



# Progressive delivery strategies

## Blue/Green Deployments

### ➤ Advantages

- Enables **zero-downtime releases**
- Supports **quick rollback** in case of errors
- Simplifies testing under production-like conditions

### ➤ Considerations

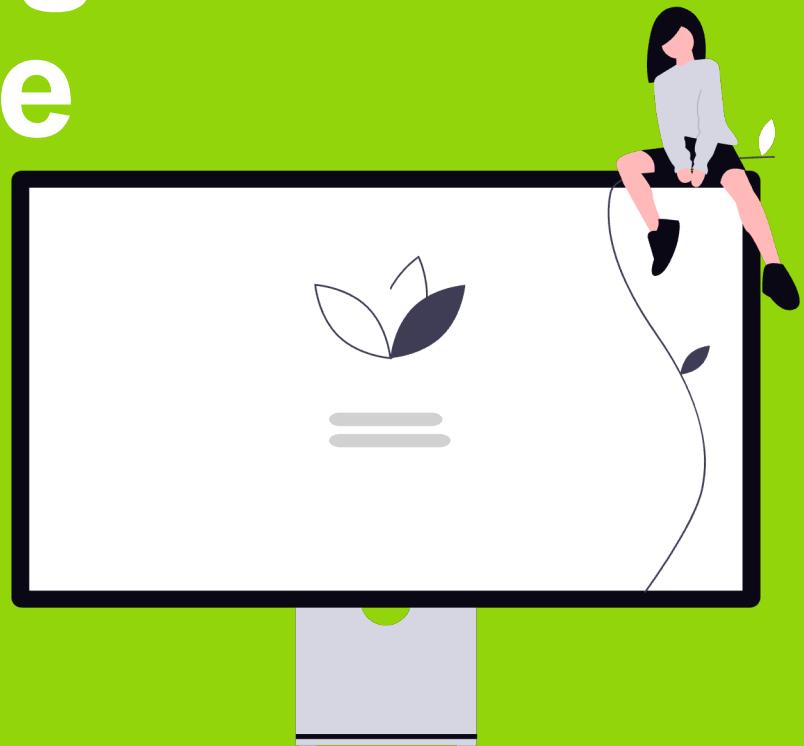
- Databases are shared or duplicated across environments
- Switching requires data migration or synchronization
- Consistency and latency must be carefully managed

# Progressive delivery strategies

## Relationship between strategies

Strategy	Key Focus	Rollback Type	User Exposure
Progressive Rollout	Gradual exposure	Partial	Controlled increase
Canary Release	Early validation	Partial	Selected users
Blue-Green Deployment	Full environment swap	Instant	Complete switch
A/B Testing	Behavioural comparison	N/A	Parallel testing groups

# Continuously releasing user-installed software



# Continuously user software releases

- Releasing a new version of your application to a production environment you control is one thing.
- Releasing a new version of software installed by users on their own machines (client-installed software) is another.
- There are several issues to consider:
  - Managing the upgrade experience
  - Migrating binaries, data, and configuration
  - Testing the upgrade process
  - Getting crash reports from users

# Continuously user software releases

## Problem

- Multiple versions of software “in the wild” cause support headaches.
- Debugging requires reverting to specific historical versions.
- Ideally: all users run the latest stable version.
- Goal: make upgrades painless and seamless.

# Continuously user software releases

## Upgrade strategies (client)

- Check for new version, prompt user to download & upgrade.
  - Easy to implement, frustrating for users.
- Download in background, prompt for installation.
  - Less intrusive, but repeated prompts can annoy.
- Download silently, upgrade on restart (Firefox model).
  - Most user-friendly if done correctly.

# Continuously user software releases

## Pitfalls of user choice

- Users often **delay upgrades** due to fear of breaking software.
- Providing choice can signal lack of developer confidence.
- Rational user decision: avoid upgrading if process seems risky.

# Continuously user software releases

## Correct approach

- Make upgrades bulletproof and silent.
- Automatically revert to previous version if upgrade fails.
- Report failures to development team without bothering the user.
- Only prompt users if corrective action is needed.

# Continuously user software releases

**Don't delete the old files, move them**

- Make sure you keep a copy of the previous version around
  - Then, ensure that you clear out the old files before deploying the new version.
  - If a stray file from the old deployment is still lying around in the newly deployed version, it can cause hard-to-track-down bugs.
- A good practice in the UNIX world is to **deploy each version** of the application into a **new directory** and **have a symbolic link** that points to the current version
  - Deploying and rolling back versions is simply a matter of changing the symbolic link to point to the previous version.
- The network version of this is to have different versions sitting on different servers or different port ranges on the same server.
  - Switch between them using a reverse proxy (Blue-Green deployments)

# Practical Guide



# Practical Guide

- Install a local version to learn and later deploy it on your server.
  - <https://github.com/flagsmith/flagsmith>
- Quick start
  - <https://docs.flagsmith.com/quickstart>